

Assignment: Iterative Closest Point (ICP)

Oscar Ramos (oscar.ramos@duke.edu)
Kris Hauser (kris.hauser@duke.edu)

February 6, 2015

1 Background

The Iterative Closest Point is an algorithm used to minimize the difference between two objects related by a rigid-body transformation, with the objective of recovering the most likely transformation. Given lists of corresponding points on each object:

$$\begin{aligned} S_1 &= \{p_1, p_2, \dots, p_n\} \\ S_2 &= \{r_1, r_2, \dots, r_n\} \end{aligned}$$

where $p_i, r_i \in \mathbb{R}^k$, the objective is to find a translation $t \in \mathbb{R}^k$ and rotation $R \in SO(k)$ that solve the problem

$$\min_{R, t} \sum_i d(p_i, Rr_i + t) \tag{1}$$

where $k = 2$ or 3 , and $d(\cdot, \cdot)$ is a distance such as the Euclidean distance and $SO(k)$ is the group of rotation matrices in \mathbb{R}^k .

An important first step is to estimate the corresponding pairs S_1 and S_2 . Typical matching approaches include exhaustive closest point matching, projections, space partitioning (k -d trees), amongst others. A variety of outlier rejection techniques may be used here as well. The ICP algorithm alternates the two steps of correspondence detection and error minimization until some notion of convergence is reached.

In this assignment we will start with a very simple 2D case and then we will address the problem for the 3D case.

2 ICP in 2D

In this part, we will focus on implementing the ICP algorithm for the planar case using two simple images. You should place all your work in the file `triangles_2d.py`.




Figure 1: Two triangles in 2D

Each image contains the same triangle but with a different position and orientation as shown in Fig. 1. Recall that in 2D, a translation is described by a position vector $t = (x_t, y_t) \in \mathbb{R}^2$, and a rotation by a matrix $R \in SO(2)$ with components

$$R = \begin{bmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{bmatrix}$$

where θ is the angle of rotation. The objective is to find the translation t and the rotation R of the second triangle with respect to the first triangle.

2.1 Preprocessing

To load images using Python, the Python Imaging Library (PIL) needs to be used. and it can be installed from the developer website (<http://www.pythonware.com/products/pil/>). The images of the triangles to be used in this part can be found in the folder `data/2d`. Note that, like in real images, there is some noise around the triangles. Before starting ICP, it is a good idea to remove that noise. The script will load the files and extract only those points whose values are less than 128 (darker than 50% grey).

2.2 Finding correspondences

In this exercise we are dealing with a small image and there are not many points (less than 200 points) that constitute the triangles. Therefore, we can apply an exhaustive search to find the closest points. (However, note that this approach is not acceptable in cases where there are many points such as for the 3D point clouds in the next section.)

Task 1. Using the square of the Euclidean distance

$$d(p_i, r_i) = \|p_i - r_i\|^2$$

to measure how close the points p_i and r_i are, find the matching pairs for the two triangles using an exhaustive search. In other terms, for every point in triangle 1, find the closest point in triangle 2. Your code for this question should go in `get_closest_points_2d()`.

Usually the result from finding the closest points contains some (or many) outliers which need to be rejected since otherwise they would contaminate the data when computing the following steps. Thus, it is a good idea to apply a threshold on the matching points based on how close they are (using the chosen distance measure). This can be done after finding the closest points or it can be done at the moment of finding them.

Task 2. Implement an outlier rejection criterion in `threshold_closest_points()` that discards the candidate pairs with distance exceeding some threshold. Later you will play around with the `threshold` value in `get_correspondences()` to make your algorithm more robust.

2.3 Minimizing the error

After the matching pairs between the two sets (triangles) has been found, the next step is to find the rotation R and translation t that minimizes the error function. In this case, using the Euclidean distance as the error function, (1) will be written as

$$\min_{R,t} \sum_{i=1}^N \|p_i - (Rr_i + t)\|^2 \quad (2)$$

where p_i and r_i are the corresponding points in the first and second triangles, respectively, and we assume that there are N correspondences.

Task 3. Implement a method that minimizes (2) in `icp_step()`. This minimization problem can be solved in different ways. One of them is using an SVD decomposition [1]. You are suggested to use the SVD solution, since it is relatively straightforward. However, if you prefer, you can implement any other method that solves (1), such as gradient descent. To implement the SVD method, first the centroids for both triangles need to be computed. They are given by

$$\mu_p = \frac{1}{N} \sum_{i=1}^N p_i \quad \mu_r = \frac{1}{N} \sum_{i=1}^N r_i$$

and then these centroids can be subtracted from the initial points

$$\tilde{p}_i = p_i - \mu_p \quad \tilde{r}_i = r_i - \mu_r.$$

It can be shown [1] that using the previous operations and replacing into the original problem, the minimization problem (1) is reduced to

$$\min_R \sum_{i=1}^N \|\tilde{p}_i - R\tilde{r}_i\|^2. \quad (3)$$

The covariance matrix of the points corresponding to both triangles can be computed as

$$H = \sum_{i=1}^N \tilde{p}_i \tilde{r}_i^T = \begin{bmatrix} \sum_{i=1}^N \tilde{p}_{ix} \tilde{r}_{ix} & \sum_{i=1}^N \tilde{p}_{ix} \tilde{r}_{iy} \\ \sum_{i=1}^N \tilde{p}_{iy} \tilde{r}_{ix} & \sum_{i=1}^N \tilde{p}_{iy} \tilde{r}_{iy} \end{bmatrix}$$

where the components of the 2D points are $\tilde{p}_i = (\tilde{p}_{ix}, \tilde{p}_{iy})$ and $\tilde{r}_i = (\tilde{r}_{ix}, \tilde{r}_{iy})$. Using an Singular Value Decomposition (SVD), the covariance matrix H can be rewritten as

$$H = U \Sigma V^T$$

where U and V are orthogonal matrices and Σ is a diagonal matrix containing the singular values of H in its diagonal. It can be shown [1] that using the SVD decomposition of the covariance matrix, the solution to (3) is given by

$$R = UV^T \quad \text{and} \quad t = \mu_p - R\mu_r. \quad (4)$$

Actually, there is another potential solution that is a critical point of the minimization, which is $R = VU^T$. One of these is a minimum and the other is a maximum. Your algorithm should test for which one is a minimizer. To solve the SVD you should use the `numpy.linalg.svd(H)` method from the numpy package (the scipy package is acceptable as well). Numpy is included in most Python distributions. You can put your code for this in `compute_error_minimizing_rotation()`.

Task 4. Debug your ICP algorithm and play with the initial conditions and the thresholding parameter. What is the best value for the parameter? Why do you think this is the case? Note: if you have the matplotlib Python package installed, the triangles and their correspondences at each stage of the ICP will be displayed to you. This should be quite helpful for debugging.

3 ICP for a 3D point cloud

In this part of the assignment we will extend the usage of ICP to the 3D case. For the amazon picking challenge, several point clouds have been acquired for different objects and they are available at http://r11.berkeley.edu/amazon_picking_challenge. The point clouds that we will be using originally come from that database, but their format has been modified to be easily read in Python.

3.1 Matching with a Known Object

The first question will ask you to implement ICP for 3D point clouds. The models of some selected objects in the APC can be found in `data/models` and are in json format. The depth maps to be used in this part can be found in `data/processed_depth`, also in json format. To read these files, use the script called `match_3d.py` which will read the `school_glue` file by default. OpenGL will show the model in its reconstructed colors, the depth map in blue, and an axis at the coordinates origin (red, green and blue represent the x , y and z axis, respectively). For testing, you should use any of the four depth maps in `data/processed_depth`.

Task 5. Implement the `icp()` function in `match_3d.py`. The steps to follow for the implementation are similar to the 2D case. However, there are several issues you must address when scaling up to 3D.

First, there are many points in the depth map and an exhaustive search is not a practical approach (you can try, but it will take a lot of time). One of the simplest approaches to addressing this challenge is to sample the data (both data sets) uniformly or randomly before attempting to match the points. Another approach can be to use a faster point location data structure, such as a grid or a K-d tree.

Also, note that there are some points in the depth map that are very close to each other within some small margin. It can be useful to only keep one of those points before doing the pair matching in order to have fewer possible pairs. Another problem that might occur is that some points can be too close to the model and can generate a considerable bias. In this case, you might enforce a limit on the number of points that may be matched to a single point.

The translation and rotation matching portion of this task is very similar to that of the 2D case. If the SVD method is used, note that the covariance matrix H from the 2D case needs to be modified for the 3D points as

$$H = \sum_{i=1}^N \tilde{p}_i \tilde{r}_i^T = \begin{bmatrix} \sum_{i=1}^N \tilde{p}_{ix} \tilde{r}_{ix} & \sum_{i=1}^N \tilde{p}_{ix} \tilde{r}_{iy} & \sum_{i=1}^N \tilde{p}_{ix} \tilde{r}_{iz} \\ \sum_{i=1}^N \tilde{p}_{iy} \tilde{r}_{ix} & \sum_{i=1}^N \tilde{p}_{iy} \tilde{r}_{iy} & \sum_{i=1}^N \tilde{p}_{iy} \tilde{r}_{iz} \\ \sum_{i=1}^N \tilde{p}_{iz} \tilde{r}_{ix} & \sum_{i=1}^N \tilde{p}_{iz} \tilde{r}_{iy} & \sum_{i=1}^N \tilde{p}_{iz} \tilde{r}_{iz} \end{bmatrix}$$

but the general procedure remains similar.

3.2 Matching unknown objects

Your final task will be to use ICP to attempt to discover which object is present in an “unknown” scene. The `detect_3d.py` file will run your ICP algorithm from Task 5 on all four known objects to match them to a given scene. You will also implement a scoring function that will assess the matching error for a given ICP solution. This function should help you distinguish between objects in a given scene.

As an error function you may wish to simply return the sum of squared distances in the ICP correspondences, or do something different to achieve better performance.

Task 6. Implement a `matching_error(object,scene)` function in `detect_3d.py` that (hopefully) gives lower errors to objects that are matched well to a given scene. Investigate its performance empirically in detecting the correct object on the four scenes (you will do this by changing the `scene` variable). Report your results in the commented-out area provided for you.

References

- [1] Arun, K.S., Huang, T.S. and Blostein S.D., Least-squares fitting of two 3-D point sets, *IEEE Transactions on Pattern Analysis and Machine Intelligence*. 9(5): 698-700, 1987.